

1.

- a. 100
- b. Since the addition can be pipelined, at steady state, a functional unit can process 1 addition each cycle. And since there are 4 functional units, the peak performance is 4 additions per cycle.
- c. Even though the addition can be pipelined and there are a total of 4 functional units, since we are performing a dependent chain of computation, we can't really take advantage of them. As the result, it will take 200 cycles to complete the sum.
- d. The time it takes to compute the vector sum was calculated to be 200 cycles. Since the peak performance was calculated to be 4 additions per cycle, theoretically 800 independent float point additions can be performed. The code above only has a efficiency of 0.5 addition per cycle, so it only has 12.5% of the peak performance.
- e. Since the latency is 2 cycle, 2 chains of independent addition can fill in the pipeline. And since we have 4 functional units, we need a total of 8 independent additions.

2.

- a. Since we can only fire one instruction at a cycle, there are a total of 2 operations, and each operation takes one cycle, it takes 2 cycles to complete one max. Therefore 0.5 max per cycle.
- b. Since the functional unit is pipelined, we can fill the pipeline with 2 independent chain of instructions. Then, we can perform 2 max in 4 cycles, with is still 0.5 max per cycle.
- c. For each max operation, we need to perform 2 instructions: a GEQ and a BLEND. The GEQ instruction takes in 2 inputs and produces 1 output, the BLEND instructions takes in 3 inputs and produces 1 output, but the output of BLEND can safely overwrite the output of GEQ, so the 1 max operation uses 3 registers. Furthermore, since the code is always comparing each element in the in array with 0, the subsequent max only needs 1 input. With 16 registers, we can perform 7 max operations at the same time (uses 15 registers). But since we only have 1 functional unit, it only takes 2 independent chains of max operation to fill the pipeline. The pseudocode for the kernel is then:


```
For i between 0 and len:
    input_1 = in[i]
    input_2 = in[i + 1]
    mask_1 = GEQ(input_1, 0)
    mask_2 = GEQ(input_2, 0)
    out[i] = BLEND(input_1, 0, mask_1)
    out[i + 1] = BLEND(input_2, 0, mask_1)
```

This assumes the length of the in array is divisible by 2.

3.

- a. If the FMA takes 6 cycles, then we need at least 6 independent chains of instructions to fill the pipeline. This means the matrix C must have at least 6 entries, so $M \times N$ could be either 3×2 or 6×1 .
- b. With 4 functional units, we need 24 chains of instructions to fill the pipeline. For a matrix C with 24 entries, it can have shape of 6×4 , 8×3 , 12×2 , or 24×1 .
- c. With a SIMD FMA that has a latency of 5 cycles and throughput of 2 instructions per cycle, we need at least a total of 10 independent instructions to fill the pipeline. The size of C however, depends on how many inputs 1 FMA can take. For a FMA that takes 4 inputs, C must have at least 40 entries. For the simplicity of SIMD instructions, we can let C have the shape 10×4 .

Operation/cycle vs. k



4. Around 5 hours.